

Review 5 M2 ERTS PROJECT

Validation of EDF & Implementation of ED-H on Xenomai

Our team



Debuss Alexy

C dev.



Rayella Niranjana

C dev.

- **1 Context**
- **2 Key question**
- **3 Approach**
- **4 Objectives**
- **5 Planning**
- **6 Progress**
- **7 Project evolutions**
- **8 Presentation of sprint 4**

1. Context

Meeting with the professor (M. Queudet) on october the 15th :

The interest here is to put into applications a dynamic-priority energy aware scheduling strategie ED-H through the use of a Real-time kernel.

We will use Xenomai, I mean a dual kernel configurations : a Linux kernel (using Ubuntu distributions) supplementing by a RT co-kernel (Cobalt)

A Master student recently integrated EDF into Xenomai...

2. Key question

We need to have ED-H working without affecting integrity of other scheduling policies (fixed priority, EDF).

How could we manage to integrate/validate ED-H on Xenomai ?



3. Approach

- 1) Learn about Xenomai
- 2) Validate/Assert EDF implementation on Xenomai
- 3) Understand EDF implementation
- 4) Learn about ED-H
- 5) Implement ED-H In Xenomai
- 6) Validate/Assert ED-H implementation on Xenomai

The project will be split into sprints (1 sprint = 2 weeks)

4. Primary objectives

Primary objectives are established at the beginning of the project

1) Installation of modified Xenomai and validation for EDF integration

Metric : report for modifications added to have a functional installation + tests to validate xenomai-EDF (report + source code)

Estimated time : 10h

2) Development of a linux module to grab battery information and transfer to Cobalt

Metric : guideline for the creation/use of the linux module + source code for the linux module

Estimated time : 25h

3) ED-H implementation and validation on Xenomai

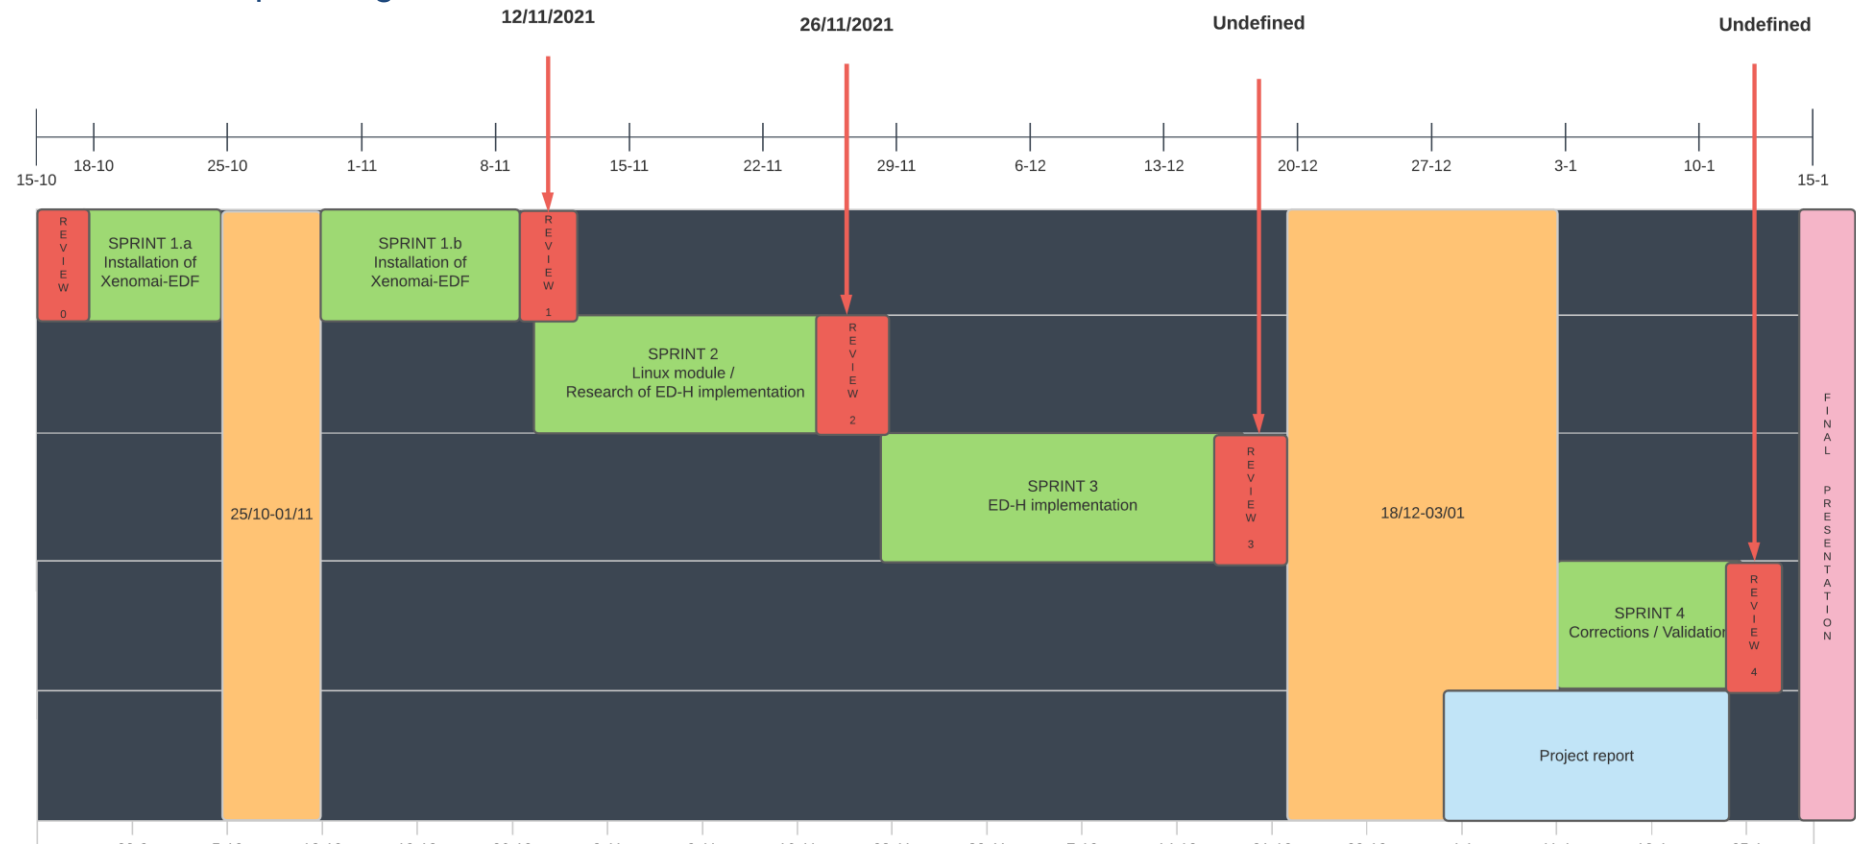
Metric : guideline for the creation/use of the linux module + source code for the linux module + source code for ED-H implementation + source code for fake battery module

Estimated time : 50h

Remaining time : 15h (Redaction of the final report, PWP presentation, handle git repo.)

5. Planning

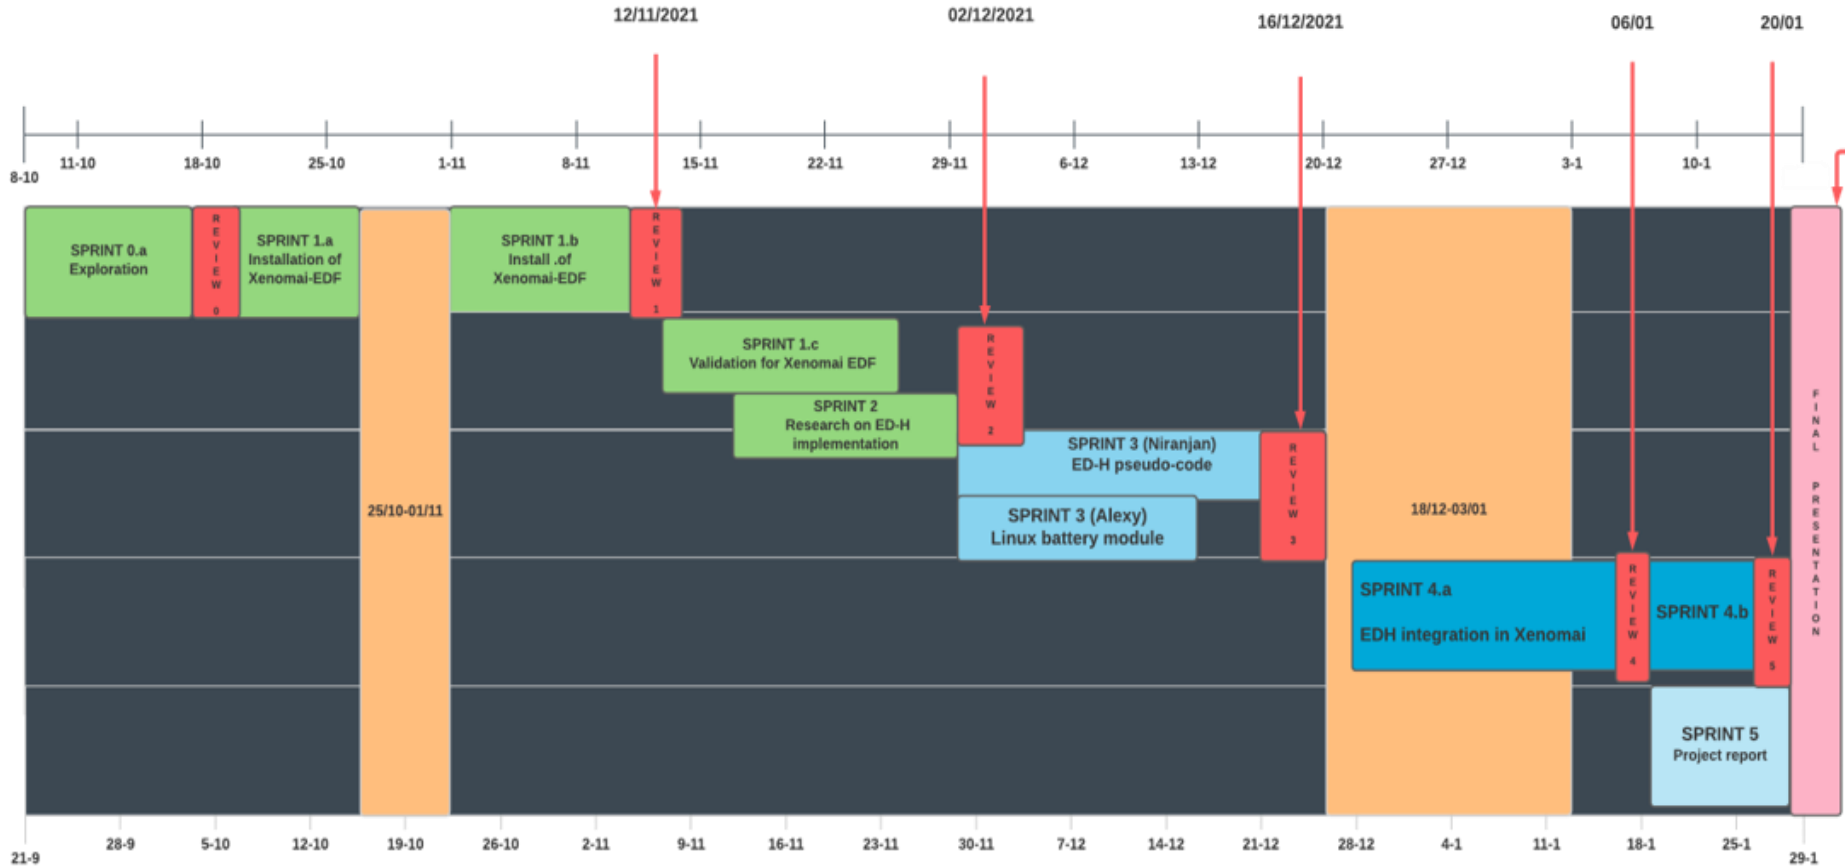
Provisional planning:





5. Planning

Effective planning:



0) Development of linux module to grab battery information

Results: Start implementation of a simple linux module (using kernel header file linux/power_supply)

2 weeks (08/10 - 17/10)

- In charge : Debus Alexy

1) Installation of Xenomai EDF

Results: Successfully installed Xenomai and validate EDF for no-preemption test cases

3 weeks (18/10 - 11/11)

- In charge : Debus Alexy & Rayella Niranjana



2) Validation of Xenomai EDF

Results: Validate EDF for preemption test cases

3 weeks (13/11 - 26/11)

- In charge : Debus Alexy



3) Linux battery module

Results: Working linux module to receive battery information

2 weeks (29/11 - 13/12)

- In charge : Debus Alexy



4) ED-H source code integration in Xenomai + battery module deployment

Results: source code for ED-H on Xenomai

(+ communication with battery module and/or fake battery module)



4 weeks ++ (17/12 - 25/01)

- In charge : Debus Alexy




Actual : (06/01 - 25/01) ED-H understanding

1. ED-H logic part.
2. End of project report

Future (25/01 – ...) :

End report redaction / Add all files/documentations in Github repo.



Actual : (06/01 - 25/01)

1. Modification for Alchemy API
2. Correction of batt.c source file for XDDP communication
3. Compilation (Kernel linux / Xenomai source code)
4. Report redaction
5. Add patch for EDF / ED-H

Future (25/01 – ...) :

Test XDDP comm. + end project documentation

8. Presentation of sprint 4.b

Modification for battery source file

```
int batt_init (void){
    //Alternative: Access to /dev/xeno_rtipc
    struct sockaddr_ipc saddr;
    int ret;
    size_t poolsz;

    /*
    * Get a datagram socket to bind to the RT endpoint. Each
    * endpoint is represented by a port number within the XDDP
    * protocol namespace.
    */

    ufd = __rtdm_dev_socket(AF_RTIPC, SOCK_DGRAM, IPCPROTO_XDDP);
    if (ufd < 0) {
        printk("XENO_WARNING __rtdm_dev_socket failed\n");
        return -1;
    }
    /*
    * Set a local 16k pool for the RT endpoint. Memory needed to
    * convey datagrams will be pulled from this pool, instead of
    * Xenomai's system pool.
    */
    poolsz = 16384; /* bytes */
    ret = rtdm_setsockopt(ufd, SOL_XDDP, XDDP_POOLSZ,
        &poolsz, sizeof(poolsz));
    if (ret)
        printk("XENO_WARNING setsockopt failed\n");
    /*
    * Bind the socket to the port, to setup a proxy to channel
    * traffic to/from the Linux domain.
    *
    * saddr.sipc_port specifies the port number to use.
    */
    memset(&saddr, 0, sizeof(saddr));
    saddr.sipc_family = AF_RTIPC;
    saddr.sipc_port = XDDP_PORT;
    ret = rtdm_bind(ufd, (struct sockaddr *)&saddr, sizeof(saddr));

    if (ret < 0)
    {
        printk(XENO_INFO "bind error\n");
        return 1;
    }
    else{
        printk(XENO_INFO "bind OK on port %d\n", XDDP_PORT);
    }

    return 0;
}
```

```
Msg_battery battery_read_msg(void)
{
    char buf[128];
    struct timespec ts;
    fd_set readfds;
    Msg_battery battery_message= {};
    int ret = 0, ret_select = 0;

    #if 0
        fcntl(ufd, F_SETFL, O_NONBLOCK);

        //TODO Warning : blocking call here... Need to use select
        ts.tv_sec = 0;
        ts.tv_nsec = 0; /* 0 ms */

        ret_select = select(ufd + 1, &readfds, NULL, NULL, &ts);

        if (ret_select == 1)
            ret = rtdm_rcvfrom(ufd, buf, MAX_BATT_READ_MSG_LENGTH, 0, NULL, 0);
    #endif

    ret = rtdm_read(ufd, buf, MAX_BATT_READ_MSG_LENGTH);
    if (ret > MAX_BATT_READ_MSG_LENGTH){
        battery_message.message_integrity = false;
    }
    if (ret >= 0 )
    {
        //Process message in struct
        sscanf((const char*)buf, "[%d,%d,%d,%d,%d]",
            (int*)&battery_message.capacity,
            (int*)&battery_message.chargenow,
            (int*)&battery_message.chargefull,
            (int*)&battery_message.battery_size,
            (int*)&battery_message.energy_production
        );

        if (battery_message.capacity >= 0 &&
            battery_message.chargenow >= 0 && battery_message.chargenow <= 100 &&
            battery_message.chargenow >= 0 && battery_message.chargenow <= 100 &&
            battery_message.battery_size >= 0 && battery_message.chargenow <= BATTERY_SIZE_MAX_VALUE &&
            battery_message.energy_production >= 0 && battery_message.energy_production <= EP_MAX_VALUE
        )
        {
            battery_message.message_integrity = true;
        }
    }
    else{
        //Handle error
        battery_message.message_integrity = false;
    }
    return battery_message;
}
```



8. Presentation of sprint 4.b

Modification for Alchemy API

```
CURRENT_DECL(int, rt_task_create(RT_TASK *task,
                                const char *name,
                                int stksize,
                                int prio,
                                int mode));

CURRENT_DECL(int, rt_task_create_dyna(RT_TASK *task,
                                     const char *name,
                                     int stksize,
                                     xnticks_t next_deadline,
                                     double WCET,
                                     double WCEC,
                                     dyna_policy policy,
                                     int mode));

CURRENT_DECL(int, rt_task_spawn(RT_TASK *task, const char *name,
                                int stksize, int prio, int mode,
                                void (*entry)(void *arg),
                                void *arg));

CURRENT_DECL(int, rt_task_spawn_dyna(RT_TASK *task, const char *name,
                                     int stksize, xnticks_t next_deadline, double WCET, double WCEC, dyna_policy policy, int mode,
                                     void (*entry)(void *arg),
                                     void *arg));
```

```
typedef enum{
    EDF = 0,
    EDH_ASAP, //As Soon As Possible
    EDH_ALAP  //As Late As Possible
}dyna_policy;
```

RTIME

8. Presentation of sprint 4.b

Function xnsched_pick_next in kernel/cobalt/sched.c

```

if (use_EDH){
    thread = &sched->rootcb;

    //ED-H Rule n°3:
    if(total_energy == 0 || slack_energy == 0){
        //Proc. IDLE
    }

    //ED-H Rule n°4:
    if(total_energy == my_msg_battery.capacity || slack_time == 0){
        thread = xnsched_rt_pick(sched);

        if (unlikely(thread == NULL))
            thread = &sched->rootcb;
    }

    //ED-H Rule n°5
    if(total_energy > 0 && total_energy < my_msg_battery.capacity
        && slack_time > 0 && slack_energy > 0){
        //Either Busy or standby : actually will be busy
    }
}
else{
    thread = xnsched_rt_pick(sched);
    if (unlikely(thread == NULL))
        thread = &sched->rootcb;
}

set_thread_running(sched, thread);

return thread;

```

- **Rule 1:** The priorities assigned by EDF are used to select the next ready job in $L_r(t_c)$.
- **Rule 2:** The processor is on standby during $[t_c, t_c + 1)$ if $L_r(t_c) = \emptyset$.
- **Rule 3:** The processor is on standby during $[t_c, t_c + 1)$ if $L_r(t_c) \neq \emptyset$ and one of the following conditions is true:
 1. $E(t_c) \approx 0$
 2. $PSE_{\mathcal{T}}(t_c) \approx 0$
- **Rule 4:** The processor is busy during $[t_c, t_c + 1)$ if $L_r(t_c) \neq \emptyset$ and one of the following conditions is true:
 1. $E(t_c) \approx C$
 2. $ST_{\mathcal{T}}(t_c) = 0$
- **Rule 5:** The processor can be either busy or on standby if $L_r(t_c) \neq \emptyset$, $0 < E(t_c) < C$, $ST_{\mathcal{T}}(t_c) > 0$ and $PSE_{\mathcal{T}}(t_c) > 0$.

Function xnsched_pick_next in kernel/cobalt/sched.c

```

struct list_head *q = &sched->rt.runnable;
struct xnthread *b_thread;
union xnsched_policy_param param;

if (list_empty(q))
    goto no_battery;

list_for_each_entry(b_thread, q, rlink) {
    if (unlikely(b_thread->sched_class == &xnsched_class_dyna)){
        //Search through EDF thread
        b_thread->sched_class->sched_getparam(b_thread, &param);

        if (param.rt.policy == EDH_ASAP || param.rt.policy == EDH_ALAP){
            use_EDH = true;
            goto battery;
        }
    }
}

/*****
battery: //Consider only one battery
// Access battery data only if one of following task use EDH scheduling class
if(use_EDH){
    my_msg_battery = battery_read_msg();

    if (my_msg_battery.message_integrity == true){
        printk(XENO_INFO
            "chargenow :%d\ncapacity:%d\n", my_msg_battery.chargenow, my_msg_battery.capacity);

        list_for_each_entry(b_thread, q, rlink) {
            if (unlikely(b_thread->sched_class == &xnsched_class_dyna)){

                /*TODO Compute for each thread:

                1) slack_time
                2) slack_energy

                using : */
                param.rt.WCET;
                param.rt.WCEC;

                my_msg_battery.chargenow;
                my_msg_battery.battery_size;
                my_msg_battery.energy_production;
            }
        }
    }
}
no_battery:

```

Compute slack_time & slack_energy for each task

***Thanks for your attention.
Do you have any questions ?***